


AgentFlow | Ultimate Technical Architecture & Interview Deep-Dive

This document is an exhaustive, ground-up technical breakdown of the **AgentFlow Platform**. It is engineered specifically for Senior/Lead AI Engineer interviews, covering exact architectural tradeoffs, system design patterns, bug resolutions, and the deep integrations that make this platform production-grade.

1. High-Level System Architecture

AgentFlow operates on a **decoupled, event-driven architecture** built on five primary layers.

1.1 Visual Representation



1.2 The Technology Stack & Justifications

- **AI Orchestration: LangGraph**
 - *Why not LangChain SequentialChains or CrewAI?* Simple chains lack cyclic capabilities. CrewAI abstracts too much state. LangGraph provides a low-level State Machine (DAG) where we maintain absolute control over the AgentState TypedDict. This allows for conditional routing, human-in-the-loop interventions, and persistent memory across complex pipelines.
- **LLM Engine: Google Gemini 1.5 & 3.1 Flash / Flash Image Preview**

- *Why?* Gemini offers a massive 1M+ token context window natively and processes multimodal tasks (text + image generation) in a single, fast SDK framework.
- **Backend: FastAPI & Uvicorn**
 - *Why?* Native `async/await` support is mandatory for I/O bound tasks like LLM API calls and WebSockets. FastAPI's Pydantic validation ensures strict schema adherence between the frontend and database.
- **Database: SQLite (via `aiosqlite`)**
 - *Why?* The project requirements demanded a "fully local, single setup command" infrastructure. SQLite requires no Docker containers or external servers. We used `aiosqlite` to prevent database locking during asynchronous concurrent agent executions.
- **Frontend: React 19, Vite, ReactFlow**
 - *Why?* ReactFlow provides an out-of-the-box, canvas-based node editor, which is essential for a visual workflow architecture tool.

2. Core Engineering Mechanisms

2.1 The LangGraph State Engine (`graph_builder.py`)

At the heart of the platform is the `AgentState` object. As execution passes from Node A to Node B, the state is appended, not overwritten.

- **Role Alternation:** Gemini strictly requires alternating `user` and `model` roles. When an agent hands over data to the next agent, the system automatically injects a `user` bridge message to maintain API compatibility and prompt the next agent properly.

2.2 Local-First Multimodal Storage (`gemini_11m.py`)

The Challenge: Initially, the image generator returned Base64 strings (often >1MB for high-res images). Passing this through the WebSocket for real-time logs caused React's Virtual DOM to choke and crash the browser. **The Engineering Fix:**

1. The backend intercepts the raw `image_bytes` .
2. Uses Python's `PIL` (Pillow) and `io.BytesIO` to write the image directly to the host's `/static/generated_images/` folder as a unique UUID `.png` .
3. The backend returns a lightweight Markdown pointer: `![Generated Image] (/static/generated_images/uuid.png)` .
4. Vite's proxy automatically routes the frontend request to FastAPI to serve the image asynchronously.

2.3 Bulletproof Tooling & Anti-Bot Bypassing (`tools.py`)

- **YouTube Transcript Extraction (`yt-dlp`)**: YouTube actively blocks automated requests (HTTP 429 / ParseErrors). We migrated from standard scrapers to `yt-dlp`, which simulates client behavior, handles internal `json3` subtitles, and bypasses almost all anti-bot walls to guarantee data retrieval.
- **Intelligent Tool Fallbacks**: If a tool fails completely, the `try/except` block returns a specific system string instructing the LLM to use a secondary tool. (e.g., *"Transcript blocked. Use web_search to find a summary instead."*)

2.4 WebSocket Telemetry & React Memoization

- **Backend (`websocket.py`)**: A `ConnectionManager` manages active sessions. Every time an LLM executes a tool, or finishes a thought, it fires an `await manager.broadcast()` event.
- **Frontend (`Monitor.jsx`)**: Receiving 10+ updates per second causes massive re-rendering. We utilized `React.memo` with a custom comparison function for the `MemoizedMarkdown` and `LogEntry` components. This means only *new* log entries render; existing DOM nodes are untouched.

2.5 Telegram Bot Concurrency & Conflict Shielding

- **The Issue**: When Uvicorn hot-reloads, it leaves orphaned Python processes polling the Telegram API, resulting in `telegram.error.Conflict` (multiple bots competing for one token).
- **The Fix**:
 - i. Hooked the bot directly into FastAPI's `@asynccontextmanager` `lifespan` events to ensure `bot.shutdown()` is called gracefully when the server terminates.
 - ii. Added HTML Markdown sanitization to prevent Telegram from crashing on unescaped special characters (e.g., `_`, `*`, `[`) emitted by the LLM.

Phase 1: Backend Foundation & Setup

This phase establishes the complete foundation for our FastAPI backend. We will create the necessary project structure, define our database schema, set up our API models, and create the initial API endpoints for managing agents and workflows.

Key Steps:

1. **Project Structure**: Create the main directories (`backend/` , `frontend/` , `docs/`) to keep the project organized from the start.
2. **Database (`database.py`)**:
 - Define the complete `SQLite` database schema in SQL. This includes tables for `agents` , `workflows` , `runs` , `messages` , and `agent_memory` .

- Implement an asynchronous connection manager (`get_db`) for interacting with the SQLite database.
- Create an `init_db` function that sets up the database and creates all tables on application startup.

3. API Models (`models.py`):

- Define all `Pydantic` models that will be used for API request and response validation. This ensures our API is type-safe and self-documenting.

4. API Routers (`routers/`):

- Create placeholder files for all our API endpoints (`agents.py` , `workflows.py` , `runs.py` , `messages.py` , `websocket.py`).
- Implement the full CRUD (Create, Read, Update, Delete) functionality for the `/api/agents` endpoint.

5. Main Application (`main.py`):

- Set up the main `FastAPI` application instance.
- Configure CORS middleware to allow the frontend application (running on a different port) to communicate with the backend.
- Implement a `lifespan` event handler to initialize the database when the server starts.
- Include all the API routers into the main application.

6. Dependencies:

- Create a `requirements.txt` file listing all necessary Python packages.
- Create a `.env.example` file to document the required environment variables.

By the end of this phase, we will have a fully functional, albeit simple, backend server that can create and manage agents, backed by a persistent SQLite database.

Phase 2: Core Agent Runtime

This phase focuses on building the "brain" of the platform. We will implement the logic that allows agents to think, use tools, and communicate, all orchestrated by LangGraph.

Key Steps:

1. **Directory Structure:** Create the `backend/runtime/` and `backend/telegram/` directories to house the core logic.
2. **Gemini LLM Wrapper (`runtime/gemini_llm.py`):**

 - Create a class to encapsulate the `google-genai` SDK.
 - This wrapper will handle communication with the Gemini API, including system prompts, temperature settings, and message history.

3. **Tool Implementation (`runtime/tools.py`):**

 - Define a set of built-in tools that agents can use (e.g., `web_search` , `calculator`).
 - Each tool will be decorated with LangChain's `@tool` decorator for easy integration.

4. **Agent Memory (`runtime/memory.py`):**

 - Implement a class that connects to the `agent_memory` table in SQLite.
 - This will provide agents with persistent, long-term memory across different runs. Agents will be able to read their memory context and save new information.

5. **Graph Builder (`runtime/graph_builder.py`):**

 - This is the core of the LangGraph implementation.
 - Create a function `build_workflow_graph` that dynamically constructs a `StateGraph` from a workflow configuration stored in the database.
 - It will create a "node" for each agent in the workflow and connect them with "edges" as defined.

6. **Workflow Executor (`runtime/executor.py`):**

 - Create the main `execute_workflow` function.
 - This function will be responsible for:
 - Loading agent and workflow configurations from the database.
 - Calling the `graph_builder` to compile the LangGraph graph.
 - Invoking the graph with the initial user input.
 - Persisting every message (from users, agents, and tools) to the `messages` table.
 - Broadcasting real-time events to the frontend via the WebSocket connection.
 - Updating the final status of the run in the `runs` table.

Phase 3: Workflow Builder & Live Monitoring

This phase is about bringing the "orchestration" part of the platform to life. We will build the visual workflow editor and the real-time monitoring dashboard.

Key Steps:

1. **Install React Flow:** Add the `reactflow` package to the frontend project. This library is the core of our visual builder.

2. **Workflow Builder UI (`pages/WorkflowBuilder.jsx`):**

- Create the main component for the "Workflows" page.
- The layout will be a two-panel design:
 - A sidebar that lists all available agents. Users can drag these agents onto the canvas.
 - A main canvas area powered by `ReactFlow`.
- Implement the logic to add agent nodes to the canvas and connect them with edges.
- Add input fields for the workflow name and a test message.
- Create a "Save & Run" button that:
 - a. Serializes the node and edge data into the required JSON format.
 - b. Calls the backend API to save the new workflow.
 - c. Triggers a new run of that workflow with the test message.

3. **Monitoring UI (`pages/Monitor.jsx`):**

- Create the component for the "Monitor" page, which will provide a live view of agent activity.
- This component will have a multi-panel layout:
 - **Run History:** A sidebar listing recent runs, their status (e.g., `running`, `completed`), and their trigger source (`manual`, `telegram`).
 - **Live Event Stream:** A panel that connects to the `/ws/logs` WebSocket and displays incoming events in real-time as they are broadcast from the executor.
 - **Message Details:** A panel that displays the full, persisted message history for a selected run from the history list. This allows users to see the complete conversation between agents and tools.

AgentFlow | Workflow Templates, Agents & Tokenomics

This document provides a comprehensive technical overview of AgentFlow's predefined workflow templates, the specialized agents powering them, and the underlying token/cost tracking architecture.

1. Workflow Templates

AgentFlow comes pre-configured with **5 production-ready workflow templates** designed for content creation, market analysis, media intelligence, and code performance auditing.

Workflow Name	Primary Objective	Agents Involved
AI Image Studio	High-fidelity image generation with structured visual prompts.	Prompt Architect, Image Generator
[TEMPLATE] Full Content Lifecycle	End-to-end content creation from research to polished article to visual assets.	Deep Researcher, Creative Writer, Visual Director
[TEMPLATE] Competitive SWAT Analysis	Market intelligence gathering and strategic competitive analysis.	Market Analyst, Quant Strategist
YouTube Insight Engine	YouTube video transcript extraction and executive summary generation.	Media Scout, Briefing Expert
Code Performance Audit	Code review focusing on logical soundness, security, and performance.	Logic Auditor, Briefing Expert

2. Agent Roster & Specializations

AgentFlow features **10 specialized agents**, each with a distinct role, system prompt, and set of tools for orchestrating complex multi-step workflows.

Visual & Design Agents

Prompt Architect (Visual Designer)

- **Role:** JSON Architect for Image Generation Models
- **Tools:** [] (none)
- **Purpose:** Analyzes content and generates structured JSON prompts optimized for image generation models. Outputs prompts with fields like `subject`, `style`, `lighting`, `composition`, and `mood`.

Image Generator (AI Artist)

- **Role:** High-Fidelity Visual Generation Specialist
- **Tools:** [] (none)
- **Purpose:** Takes structured JSON prompts and generates vivid, high-fidelity images using multimodal models. Handles image encoding and persistence.

Visual Director (Image Designer)

- **Role:** Lead Art Director
- **Tools:** ['calculator']
- **Purpose:** Interprets written content and generates high-impact visual briefs. Ensures visual output aligns with narrative themes and brand guidelines.

Research & Analysis Agents

Deep Researcher (Data Gatherer)

- **Role:** Senior Technical Researcher
- **Tools:** ['web_search', 'url_scraper']
- **Purpose:** Gathers deep technical facts, statistics, and current trends on any topic using web search and URL scraping. Focuses on credible sources and quantifiable data.

Market Analyst (Competitive Scout)

- **Role:** Competitive Intelligence Analyst
- **Tools:** ['url_scraper', 'web_search']
- **Purpose:** Deconstructs competitor strategies by scraping websites and identifying market trends, pricing models, and feature differentiation.

Quant Strategist (Financial Modeler)

- **Role:** Quantitative Financial Analyst
- **Tools:** ['calculator', 'current_time_fetcher']
- **Purpose:** Transforms market data into hard financial models, SWOT analyses, and strategic recommendations. Performs calculations and trend projections.

Content & Communication Agents

Creative Writer (Storyteller)

- **Role:** Master Content Creator

- **Tools:** [`'current_time_fetcher'`]
- **Purpose:** Takes researched facts and crafts a single, ready-to-publish article. Focuses on narrative flow, engagement, and clarity.

Media Scout (Video Intelligence)

- **Role:** AI Video Intelligence Specialist
- **Tools:** [`'youtube_transcript'`, `'url_scraper'`]
- **Purpose:** Extracts transcripts from YouTube videos using the `youtube_transcript` tool. Automatically handles blocked or unavailable transcripts with fallback strategies.

Briefing Expert (Executive Assistant)

- **Role:** Executive Assistant & Summarization Expert
- **Tools:** [`'current_time_fetcher'`]
- **Purpose:** Condenses lengthy content into clear, actionable summaries. Extracts key takeaways and delivers insights in executive-friendly formats.

Code & Quality Agents

Logic Auditor (Code Reviewer)

- **Role:** Principal Software Engineer & Security Auditor
- **Tools:** [`'calculator'`]
- **Purpose:** Analyzes code for logical soundness, security vulnerabilities (OWASP), efficiency, and clean code adherence. Provides detailed improvement recommendations.

3. Token Tracking & Cost Calculation

AgentFlow implements a precise, real-time telemetry system to monitor LLM usage and estimate operational costs across all agent executions.

3.1 Token Capture Mechanism (`runtime/gemini_llm.py`)

Token usage is captured **at the moment of execution** for every agent interaction:

Primary Method: Native Metadata

- The system reads `response.usage_metadata.total_token_count` directly from the Google Gemini SDK.

- This provides exact token consumption from the LLM provider.

Fallback Method: Heuristic Estimation

- If metadata is unavailable, usage is estimated using: `len(generated_text.split()) + 50`
- The `+ 50` padding accounts for system prompts and framework overhead.

Multimodal Logic: Image Generation

- Image generation (using `gemini-3.1-flash-image-preview` or similar) is assigned a **fixed cost of 1,000 tokens** per image.
- This reflects the higher computational overhead of visual generation compared to text.

3.2 Cost Calculation Formula (`frontend/src/pages/Monitor.jsx`)

The platform translates tokens into estimated USD using a fixed price multiplier:

Formula:

`Estimated` Cost (USD) = Total Tokens × `$0.000002`

Price Constants:

- **Price Multiplier:** `$0.000002` per token
- **Equivalent:** `$2.00` per 1,000,000 tokens
- **Reference:** Based on Google Gemini's text generation pricing tier

3.3 Persistence & Historical Tracking (`database.py`)

Total tokens for an entire workflow run are aggregated and stored in the `usage_tokens` column of the `runs` table:

Column	Type	Purpose
<code>run_id</code>	TEXT	Unique identifier for the workflow execution
<code>workflow_id</code>	TEXT	Reference to the workflow executed
<code>status</code>	TEXT	Execution state (<code>running</code> , <code>completed</code> , <code>failed</code>)
<code>trigger_source</code>	TEXT	Execution source (<code>manual</code> from Web UI or <code>telegram</code> from bot)
<code>result</code>	TEXT	Final polished output from the last agent in the graph

Column	Type	Purpose
usage_tokens	INTEGER	Total tokens consumed across all agents
created_at	TIMESTAMP	ISO timestamp of execution

3.4 Real-Time Cost Monitoring

The **Monitor** page displays live cost tracking:

- **Live Token Count:** Updated as agents execute and consume tokens
- **Estimated Cost:** Real-time USD estimate calculated using the multiplier above
- **Run History:** Historical cost analysis for all completed executions

4. Tools Available in the Ecosystem

AgentFlow provides the following tools that agents can invoke during execution:

Tool Name	Purpose	Used By
web_search	Google Search grounding for factual data retrieval	Deep Researcher, Market Analyst
url_scraper	Website content extraction and parsing	Deep Researcher, Market Analyst, Media Scout
youtube_transcript	YouTube video transcript extraction	Media Scout
current_time_fetcher	Fetches current date/time for context	Creative Writer, Quant Strategist, Briefing Expert
calculator	Mathematical computations and financial modeling	Visual Director, Quant Strategist, Logic Auditor

5. Workflow Execution Model

All workflows operate on **LangGraph**, a directed acyclic graph (DAG) architecture where:

1. **Sequential Execution:** Agents execute in order, with output from Agent N becoming input for Agent N+1.

2. **State Management:** An `AgentState` object maintains context and message history across all agents.
3. **Tool Invocation:** Each agent can call assigned tools during execution, with results automatically injected into the prompt.
4. **Token Aggregation:** Token usage from each agent is summed into the workflow's `usage_tokens` total.
5. **Result Finalization:** The output from the last agent is stored as the `result` field in the run record.

6. Key Design Decisions

Multimodal Cost Assignment

Image generation tasks are assigned a flat 1,000-token cost to simplify cost tracking while reflecting the compute intensity. This prevents massive byte strings from inflating token counts and keeps billing predictable.

Heuristic Fallback

The fallback estimation (`len().split() + 50`) ensures that even if the LLM SDK fails to return metadata, cost tracking never goes blind. This provides visibility into API failures.

Real-Time Telemetry

WebSocket broadcasts from the backend ensure that the frontend's Monitor page reflects live execution progress and cost accumulation without polling delays.